


# Symmetric-Key Cryptography

CS 161 Spring 2022 - Lecture 7

# Announcements

- Project 1 is released
  - Checkpoint is due Friday, February 4th, 11:59 PM PT
  - Final submission is due Friday, February 18th, 11:59 PM PT
- Project party today!
  - 2-5pm in the Woz
- In person office hours & project parties really are better!
  - Much better student/unit-time support from TAs from both the student and TA viewpoint



# Cryptography Roadmap

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none"><li>• One-time pads</li><li>• Block ciphers with chaining modes (e.g. AES-CBC)</li><li>• Stream ciphers</li></ul>	<ul style="list-style-type: none"><li>• RSA encryption</li><li>• ElGamal encryption</li></ul>
Integrity, Authentication	<ul style="list-style-type: none"><li>• MACs (e.g. HMAC)</li></ul>	<ul style="list-style-type: none"><li>• Digital signatures (e.g. RSA signatures)</li></ul>

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

- Key management (certificates)
- Password management

# Symmetric-Key Encryption



Textbook Chapter 6.1

# Cryptography Roadmap

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none"><li>• One-time pads</li><li>• Block ciphers with chaining modes (e.g. AES-CBC)</li><li>• Stream ciphers</li></ul>	<ul style="list-style-type: none"><li>• RSA encryption</li><li>• ElGamal encryption</li></ul>
Integrity, Authentication	<ul style="list-style-type: none"><li>• MACs (e.g. HMAC)</li></ul>	<ul style="list-style-type: none"><li>• Digital signatures (e.g. RSA signatures)</li></ul>

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)


- Key management (certificates)
- Password management

# Symmetric-Key Encryption

- The next few schemes are symmetric-key encryption schemes
  - **Encryption schemes** aim to provide *confidentiality* (but not integrity or authentication)
  - **Symmetric-key** means Alice and Bob share the same secret key that the attacker doesn't know
    - Don't worry about how Alice and Bob share the key for now
- For modern schemes, we're going to assume that messages are *bitstrings*
  - **Bitstring**: A sequence of bits (0 or 1), e.g. 11010101001001010
  - Text, images, etc. can usually be converted into bitstrings before encryption, so bitstrings are a useful abstraction. After all, everything in a computer is just a sequence of bits!

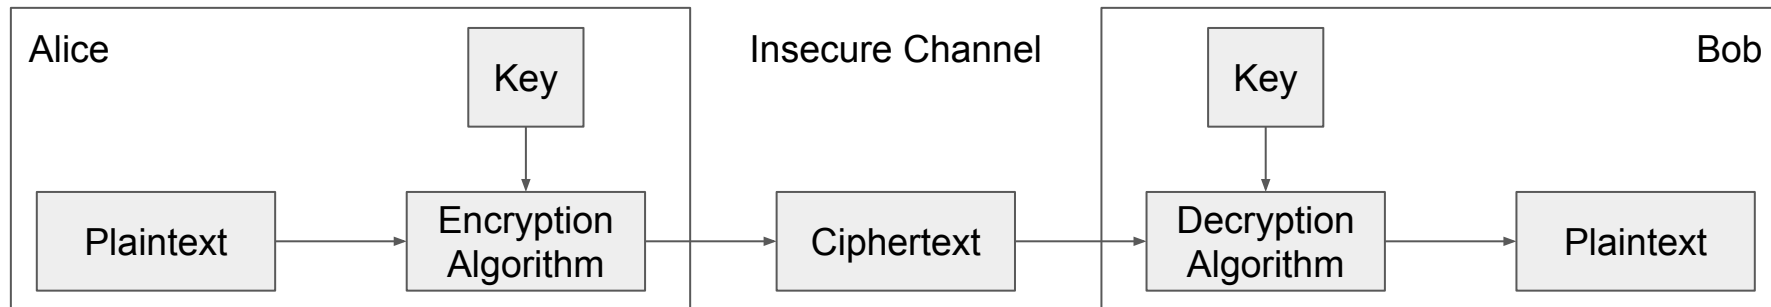
# Symmetric-Key Encryption: Definition

- A symmetric-key encryption scheme has three algorithms:
  - $\text{KeyGen}() \rightarrow K$ : Generate a key  $K$
  - $\text{Enc}(K, M) \rightarrow C$ : Encrypt a **plaintext**  $M$  using the key  $K$  to produce **ciphertext**  $C$
  - $\text{Dec}(K, C) \rightarrow M$ : Decrypt a ciphertext  $C$  using the key  $K$



# Symmetric-Key Encryption: Definition

- What properties do we want from a symmetric encryption scheme?
  - **Correctness:** Decrypting a ciphertext should result in the message that was originally encrypted
    - $\text{Dec}(K, \text{Enc}(K, M)) = M$  for all  $K \leftarrow \text{KeyGen}()$  and  $M$
  - **Efficiency:** Encryption/decryption algorithms should be fast: >1 Gbps on a standard computer
  - **Security:** Confidentiality





# Defining Confidentiality

- Recall our definition of confidentiality from earlier: “An adversary cannot read our messages”
  - This definition isn’t very specific
    - What if Eve can read the first half of Alice’s message, but not the second half?
    - What if Eve figures out that Alice’s message starts with “Dear Bob”?
  - This definition doesn’t account for prior knowledge
    - What if Eve already knew that Alice’s message ends in “Sincerely, Alice”?
    - What if Eve knows that Alice’s message is “BUY!” or “SELL” but doesn’t know which?



# Defining Confidentiality

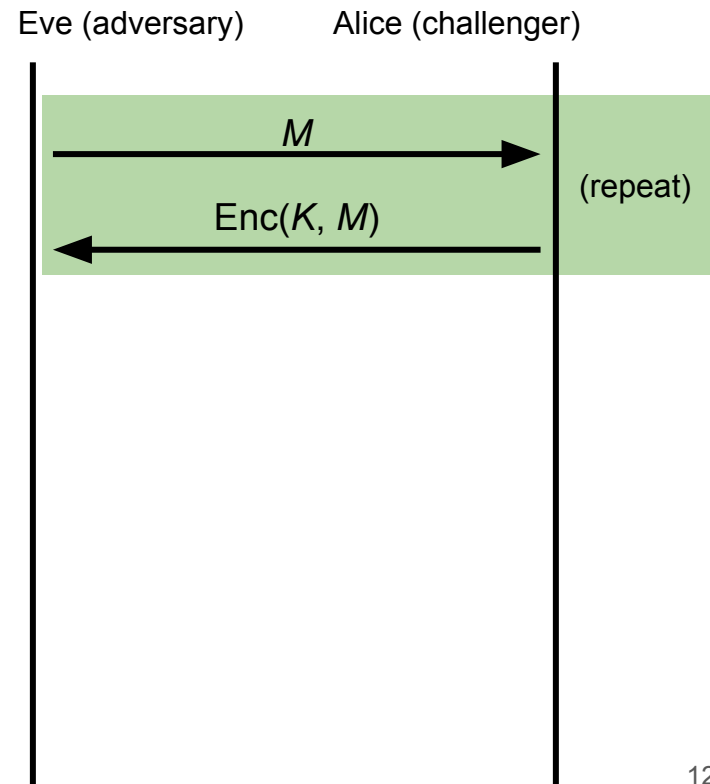
- A better definition of confidentiality: The ciphertext should not give the attacker *any additional information* about the plaintext.
- Let's design an experiment/ security game to test our definition:
  - Eve chooses two messages  $M_0$  and  $M_1$  of the same length
  - Alice chooses one message at random  $M_b$ , encrypts it, and sends the ciphertext
  - Eve knows either  $M_0$  or  $M_1$  was sent, but doesn't know which
  - Eve reads the ciphertext and tries to guess which message was sent
  - If the probability that Eve correctly guesses which message was sent is  $1/2$ , then the encryption scheme is confidential
- Intuition
  - If the scheme is confidential, Eve can only guess with probability  $1/2$ , which is no different than if Eve hadn't sent the ciphertext at all
  - In other words: the ciphertext gave Eve no *additional* information about which plaintext was sent!

# Defining Confidentiality: IND-CPA

- Recall our threat model: Eve can also perform a **chosen plaintext attack**
  - Eve can trick Alice into encrypting arbitrary messages of Eve's choice
  - We can adapt our experiment to account for this threat model
- A better definition of confidentiality: Even if Eve is able to trick Alice into encrypting messages, Eve can still only guess what message Alice sent with probability  $1/2$ .
  - This definition is called **IND-CPA** (indistinguishability under chosen plaintext attack)
- Cryptographic properties are often defined in terms of “games” that an adversary can either “win” or “lose”
  - We will use one to define confidentiality precisely

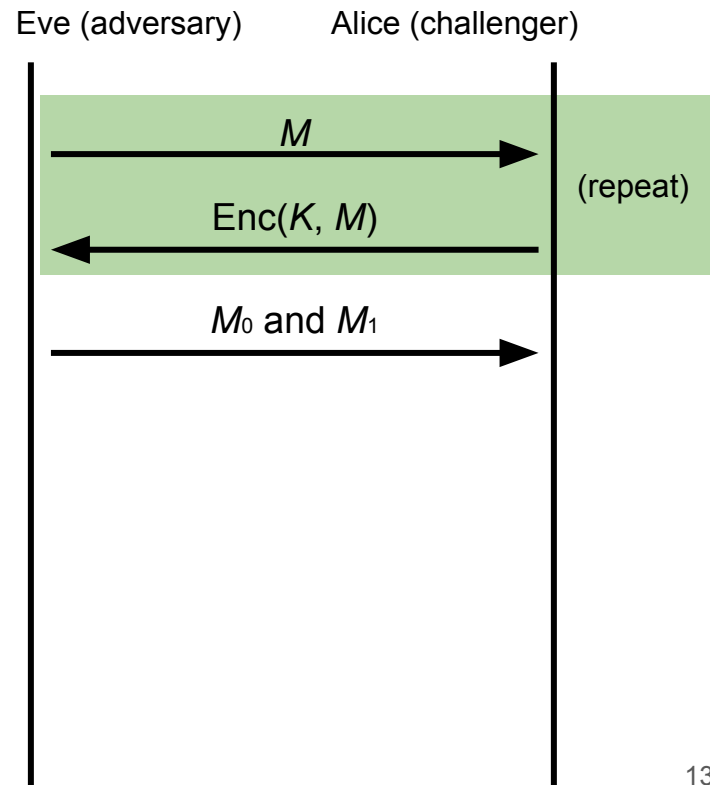
# Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts



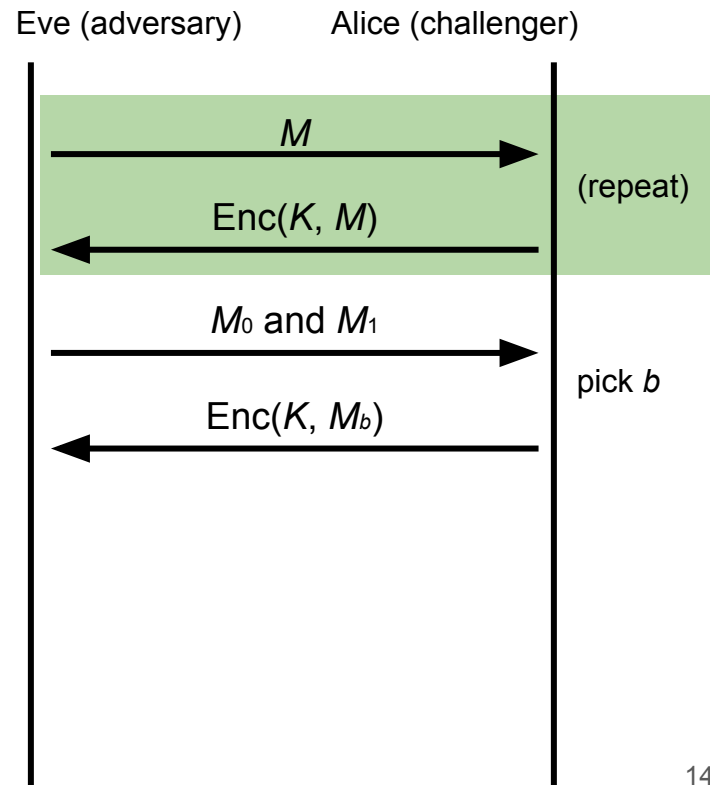
# Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts  $M_0$  and  $M_1$  to Alice



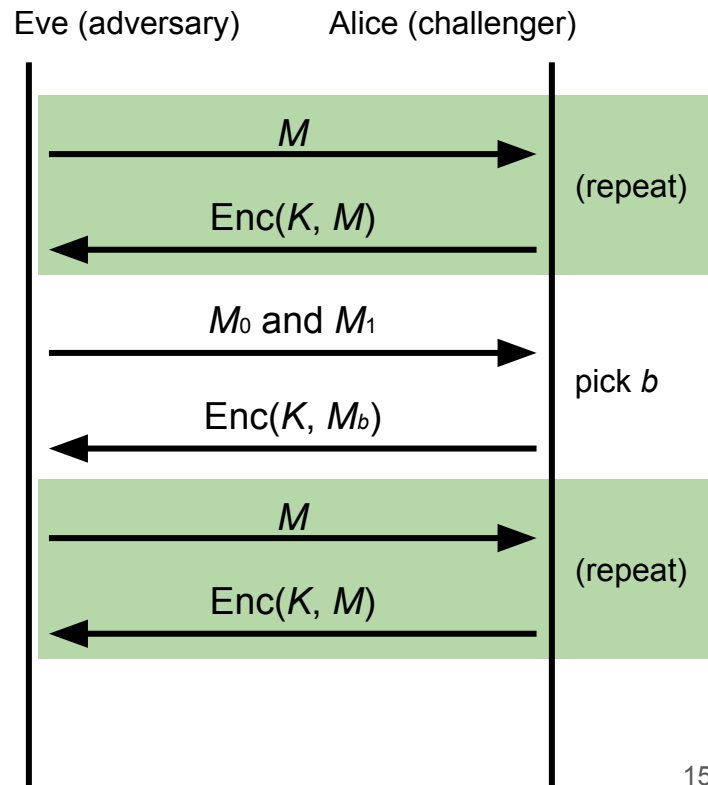
# Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts  $M_0$  and  $M_1$  to Alice
3. Alice randomly chooses either  $M_0$  or  $M_1$  to encrypt and sends the encryption back
  - a. Alice does not tell Eve which one was encrypted!



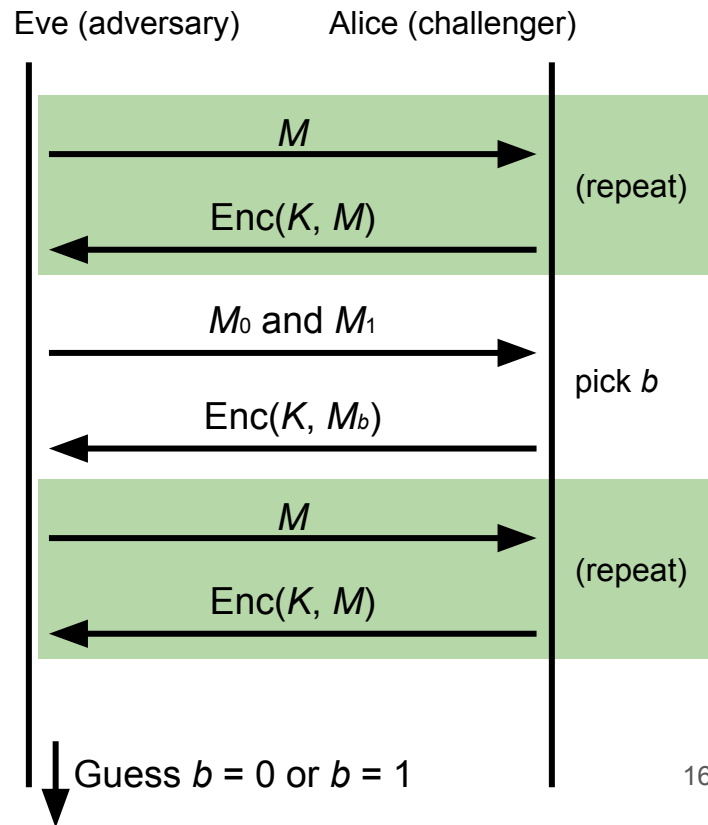
# Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts  $M_0$  and  $M_1$  to Alice
3. Alice randomly chooses either  $M_0$  or  $M_1$  to encrypt and sends the encryption back
  - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts



# Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts  $M_0$  and  $M_1$  to Alice
3. Alice randomly chooses either  $M_0$  or  $M_1$  to encrypt and sends the encryption back
  - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether Alice encrypted  $M_0$  or  $M_1$





# Defining Confidentiality: IND-CPA

- If Eve correctly guesses which message Alice encrypted, then Eve wins. Otherwise, she loses.
- How does Eve guess whether  $M_0$  or  $M_1$  was encrypted? What strategy does she use?
  - We don't *assume* she uses a particular strategy; Eve represents all possible strategies
- Proving insecurity: There exists at least *one* strategy that can win the IND-CPA game with probability  $> 1/2$ 
  - $1/2$  is the probability of winning by random guessing
  - If you can be better than random, then the ciphertext has leaked information, and Eve is able to learn it and use it to gain an advantage!
- Proving security: For *all* attackers/Eve-s, the probability of winning the IND-CPA game is at most  $1/2$

# Edge Cases: Length

- Cryptographic schemes are (usually) allowed to leak the length of the message
  - To hide length: All messages must always be the same length
    - 16-byte messages: We can't encrypt large messages (images, videos, etc.)!
    - 1-GB messages: Sending small messages (text, Tweets, etc.) needs 1 GB of bandwidth!
    - This is unpractical
  - Applications that which to hide length must choose to *pad* their own messages to the maximum possible length before encrypting
- In the IND-CPA game:  $M_0$  and  $M_1$  must be the same length
  - To break IND-CPA, Eve must learn something other than message length



# Edge Cases: Attacker Runtime

- Some schemes are theoretically vulnerable, but secure in any real-world setting
  - If an attack takes longer than the life of the solar system to complete, it probably won't happen!
  - Or if it would require a computer made out of a literal galaxy worth of science-fiction nanotech
- In the IND-CPA game: Eve is limited to a practical runtime
  - One common practical limit: Eve is limited to polynomial runtime algorithms (no exponential-time algorithms)



# Edge Cases: Negligible Advantage

- Sometimes it's possible for Eve to win with probability  $1/2 + 1/2^{128}$ 
  - This probability is greater than  $1/2$ , but it's so close to  $1/2$  that it's as good as  $1/2$ .
  - Eve's advantage is so small that she can't use it for any practical attacks
- In the IND-CPA game: The scheme is secure even if Eve can win with probability  $\leq 1/2 + \epsilon$ , where  $\epsilon$  is *negligible*
  - The actual mathematical definition of negligible is out of scope
  - Example:  $1/2 + 1/2^{128}$ : Negligible advantage
  - Example:  $2/3$ : Non-negligible advantage



# Edge Cases: Negligible Advantage

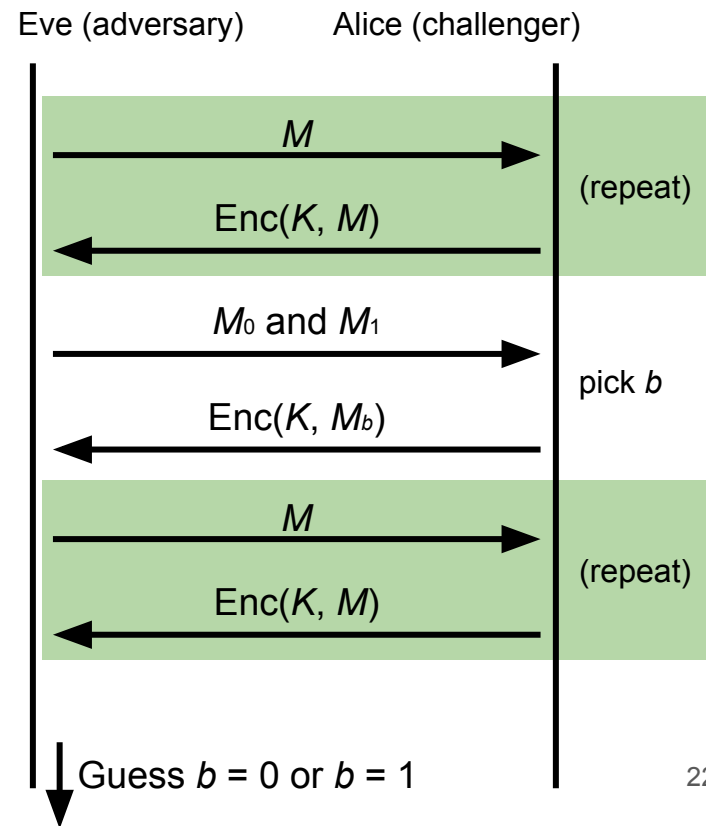
- Defining negligibility mathematically:
  - Advantage of the adversary should be exponentially small, based on the security parameters of the algorithm
  - Example: For an encryption scheme with a  $k$ -bit key, the advantage should be  $O(1/2^k)$
- Defining negligibility practically:
  - A  $1/2^{128}$  probability is completely inconceivable
  - A  $1/2^{20}$  probability is fairly likely
    - “One in a million events happen every day in New York City”
  - In between these extremes, it can be messy
    - Different algorithms run faster or slower and have their own security parameters
    - Computers get more powerful over time
    - Recall: Know your threat model!
- **Takeaway:** For now,  $2^{80}$  is a reasonable threshold, but this will change over time!

# IND-CPA: Putting it together

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts  $M_0$  and  $M_1$  to Alice
3. Alice randomly chooses either  $M_0$  or  $M_1$  to encrypt and sends the encryption back
  - Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether Alice encrypted  $M_0$  or  $M_1$

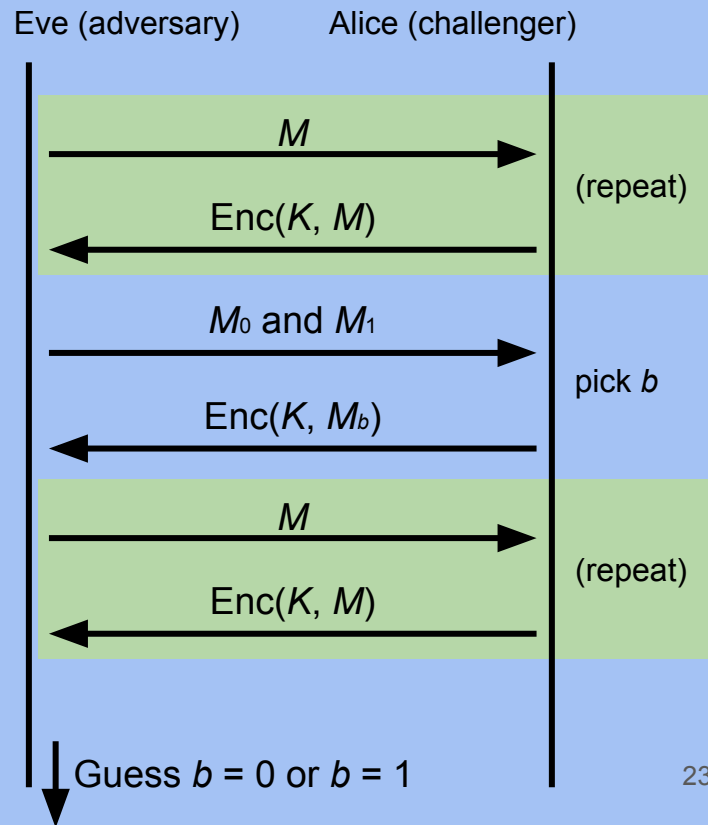


- An encryption scheme is IND-CPA secure if for all polynomial time attackers Eve:
  - Eve can win with probability  $\leq 1/2 + \epsilon$ , where  $\epsilon$  is negligible.




# Enigma: Secure under IND-CPA?

- Enigma has a significant weakness: a letter never maps to itself
  - No rotor maps a letter to itself
  - The reflector never maps a letter to itself
  - This property is necessary for Enigma's mechanical system to work
- What pair of messages should Eve send to Alice in the challenge phase?
  - Send  $M_0 = A^k$ ,  $M_1 = B^k$
  - $M_0$  is a string of  $k$  'A' characters,  $M_1$  is a string of  $k$  'B' characters
- How can Eve probably know which message Alice encrypted?
  - If there are no 'A' characters, it was  $M_0$
  - If there are no 'B' characters, it was  $M_1$



# One-Time Pads



Textbook Chapter 6.2 & 6.3



# Cryptography Roadmap

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none"><li>● One-time pads</li><li>● Block ciphers with chaining modes (e.g. AES-CBC)</li><li>● Stream ciphers</li></ul>	<ul style="list-style-type: none"><li>● RSA encryption</li><li>● ElGamal encryption</li></ul>
Integrity, Authentication	<ul style="list-style-type: none"><li>● MACs (e.g. HMAC)</li></ul>	<ul style="list-style-type: none"><li>● Digital signatures (e.g. RSA signatures)</li></ul>

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

# Review: XOR

The XOR operator takes two bits and outputs one bit:

$0 \oplus 0 = 0$
$0 \oplus 1 = 1$
$1 \oplus 0 = 1$
$1 \oplus 1 = 0$

Useful properties of XOR:

$x \oplus 0 = x$
$x \oplus x = 0$
$x \oplus y = y \oplus x$
$(x \oplus y) \oplus z = x \oplus (y \oplus z)$
$(x \oplus y) \oplus x = y$

# Review: XOR Algebra

- Algebra works on XOR too

$y \oplus 1 = 0$	Goal: Solve for $y$
$y \oplus 1 \oplus 1 = 0 \oplus 1$	XOR both sides by 1
$y = 1$	Simplify with identities

# One-Time Pads: Key Generation

Alice

$K$

0	1	1	0	0	1	0	1	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

The key  $K$  is a randomly-chosen bitstring.

Recall: We are in the symmetric-key setting, so we'll assume Alice and Bob both know this key.

# One-Time Pads: Encryption

Alice

$K$

0	1	1	0	0	1	0	1	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

$M$

1	0	0	1	1	0	0	1	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---

The plaintext  $M$  is the bitstring that Alice wants to encrypt.

Idea: Use XOR to scramble up  $M$  with the bits of  $K$ .

# One-Time Pads: Encryption

Alice

$K$	0	1	1	0	0	1	0	1	0	1	1	1
	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$
$M$	1	0	0	1	1	0	0	1	0	1	0	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
$C$	1	1	1	1	1	1	0	0	0	0	1	1

Encryption algorithm: XOR each bit of  $K$  with the matching bit in  $M$ .

The ciphertext  $C$  is the encrypted bitstring that Alice sends to Bob over the insecure channel.

# One-Time Pads: Decryption

Bob

$K$	0	1	1	0	0	1	0	1	0	1	1	1
$C$	1	1	1	1	1	1	0	0	0	0	1	1

Bob receives the ciphertext  $C$ . Bob knows the key  $K$ . How does Bob recover  $M$ ?

# One-Time Pads: Decryption

Bob

$K$	0	1	1	0	0	1	0	1	0	1	1	1
	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$	$\oplus$
$C$	1	1	1	1	1	1	0	0	0	0	1	1
	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$	$\downarrow$
$M$	1	0	0	1	1	0	0	1	0	1	0	0

Decryption algorithm: XOR each bit of  $K$  with the matching bit in  $C$ .



# One-Time Pad

- **KeyGen()**
  - Randomly generate an  $n$ -bit key, where  $n$  is the length of your message
    - Recall: For today, we assume that Alice and Bob can securely share this key
    - For one-time pads, we generate a *new* key for every message
- **$\text{Enc}(K, M) = K \oplus M$** 
  - Bitwise XOR  $M$  and  $K$  to produce  $C$ 
    - In other words: XOR the  $i$ th bit of the plaintext with the  $i$ th bit of the key.
    - $C_i = K_i \oplus M_i$
  - Alice and Bob use a different key for each encryption (this is the “one-time” in one-time pad).
- **$\text{Dec}(K, C) = K \oplus C$** 
  - Bitwise XOR  $C$  and  $K$  to produce  $M$ 
    - $M_i = K_i \oplus C_i$

# One-Time Pad: Correctness

- Correctness: If we encrypt and then decrypt, we should get the original message back

$$\text{Enc}(K, M) = K \oplus M$$

Definition of encryption

$$\text{Dec}(K, \text{Enc}(K, M)) = \text{Dec}(K, K \oplus M)$$

Decrypting the ciphertext

$$= K \oplus (K \oplus M)$$

Definition of decryption

$$= M$$

XOR property

# One-Time Pad: Security

- Recall our definition of confidentiality: The ciphertext should not give the attacker any additional information about the plaintext
- Recall our experiment to test confidentiality from earlier:
  - Alice has encrypted and sent either  $M_0$  or  $M_1$
  - Eve knows either  $M_0$  or  $M_1$  was sent, but doesn't know which
  - Eve reads the ciphertext and tries to guess which message was sent
  - If the probability that Eve correctly guesses which message was sent is  $1/2$ , then the encryption scheme is confidential

# One-Time Pad: Security

Possibility 0: Alice sends  $\text{Enc}(K, M_0)$

The ciphertext is  $C = K \oplus M_0$

Therefore,  $K = C \oplus M_0$

Possibility 1: Alice sends  $\text{Enc}(K, M_1)$

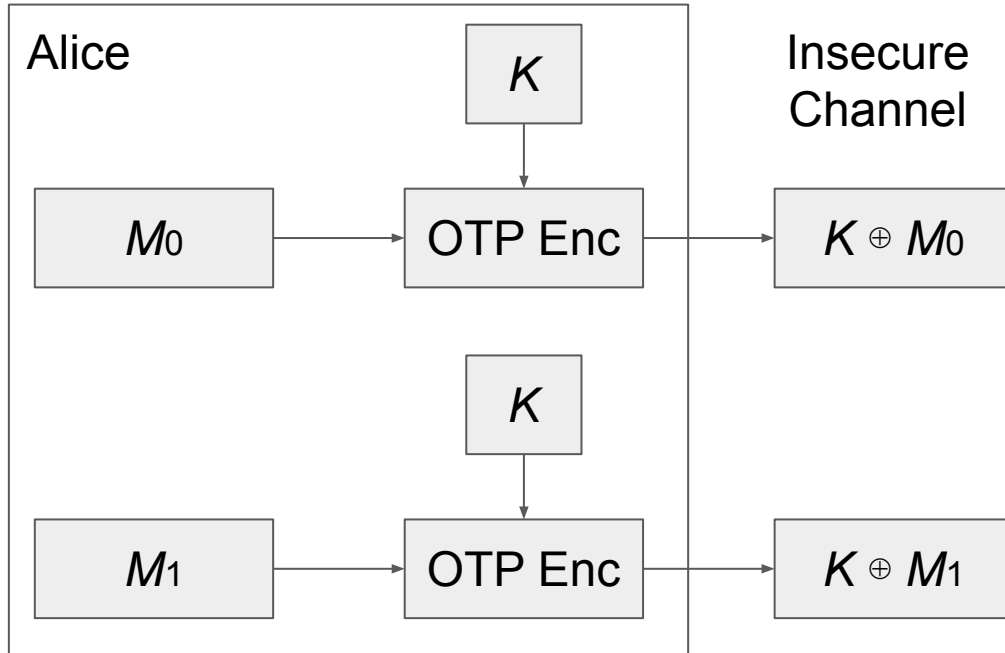
The ciphertext is  $C = K \oplus M_1$

Therefore,  $K = C \oplus M_1$

$K$  was chosen randomly, so both possibilities are equally possible

Eve has learned no new information,  
so the scheme is **perfectly secure**

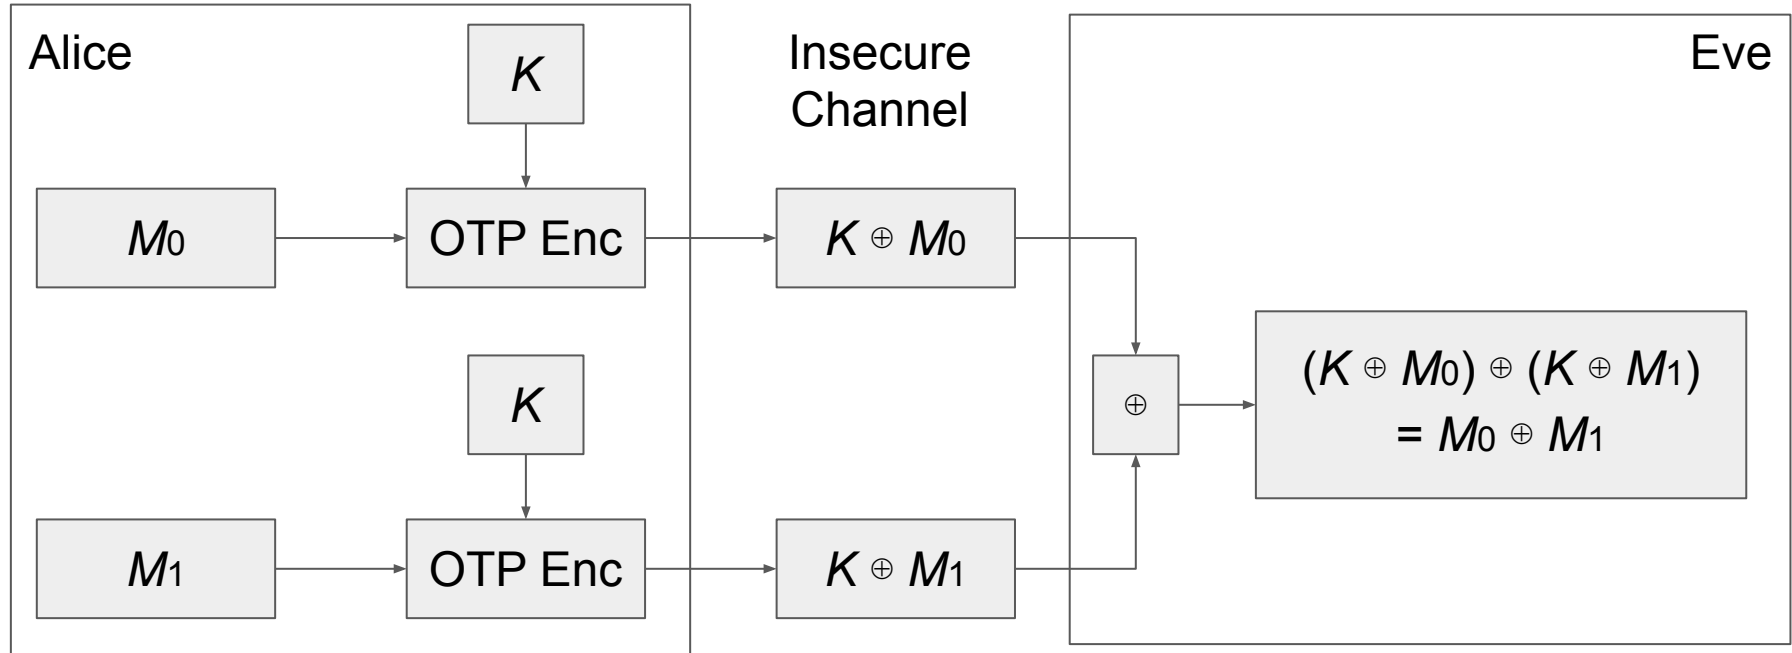
# Two-Time Pads?



Eve sees two ciphertexts over the insecure channel.

What if we use the same key  $K$  to encrypt two different messages?

# Two-Time Pads?



If Eve XORs the two ciphertexts, she learns  $M_0 \oplus M_1$ !

# Two-Time Pads?

- What if we use the same key *twice*?
  - Alice encrypts  $M_0$  and  $M_1$  with the same key
  - Eve observes  $K \oplus M_0$  and  $K \oplus M_1$
  - Eve computes  $(K \oplus M_0) \oplus (K \oplus M_1) = M_0 \oplus M_1$ 
    - Recall the XOR property: the  $K$ 's cancel out
- Eve has learned  $M_0 \oplus M_1$ . This is partial information about the messages!
  - In words, Eve knows which bits in  $M_0$  match bits in  $M_1$
  - If Eve knows  $M_0$ , she can deduce  $M_1$  (and vice-versa)
  - Eve can also guess  $M_0$  and check that  $M_1$  matches her guess for  $M_0$
- Result: One-time pads are not secure if the key is reused
  - Alice and Bob must use a different key for every message!

# Impracticality of One-Time Pads

- Problem #1: Key generation
  - For security to hold, keys must be randomly generated for every message, and never reused
  - Randomness is expensive, as we'll see later
- Problem #2: Key distribution
  - To communicate an  $n$ -bit message, we need to securely communicate an  $n$ -bit key first
  - But if we have a way to securely communicate an  $n$ -bit key, we could have communicated the message directly!
- Only practical application: Communicate keys in advance
  - You have a secure channel now, but you won't have it later
  - Use the secure channel now to communicate keys in advance
  - Use one-time pad later to communicate over the insecure channel
  - And people can compute this by hand without computers!



# One-Time Pads in Practice: Spies

- At home base, the spy obtains a large amount of key material (e.g. a book of random bits)
- In the field, the spy listens for secret messages from their home country
  - There are shortwave and terrestrial radio “number stations”
  - At a regular time, a voice gets on the air and reads a series of numbers
  - If you don’t know the key, this looks like a meaningless sequence of random numbers
  - If you know the key, you can decrypt the spy message!
- What if you don’t want to send anything to any spies?
  - Read out a list of random numbers anyway
  - Because one-time pad leaks no information, an eavesdropper can’t distinguish between an encrypted message and random numbers!

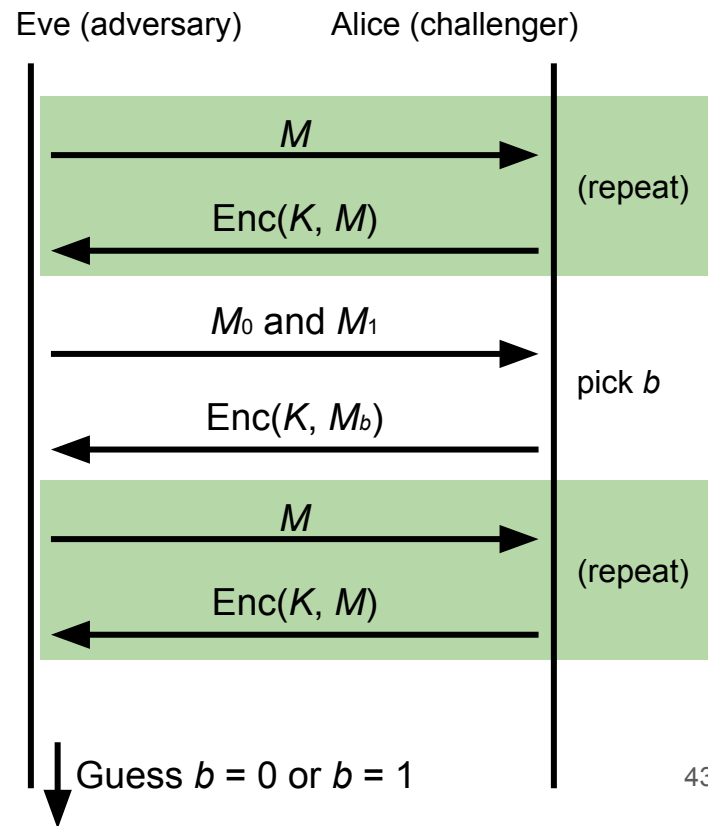
# Two-Time Pads in Practice: VENONA

- Soviet spies used one-time pads for communication from their spies in the US
- During WWII, the Soviets started reusing key material
  - Uncertain whether it was just the cost of generating pads or what...
- VENONA was a US cryptanalysis project designed to break these messages
  - Included confirming/identifying the spies targeting the US Manhattan project
  - Project continued until 1980!
- Not declassified until 1995!
  - So secret even President Truman wasn't informed about it
  - The Soviets found out about it in 1949 through their spy Ken Philby, but their one-time pad reuse was fixed after 1948 anyway
- **Takeaway:** Otherwise-secure cryptographic systems can fail very badly if used improperly!



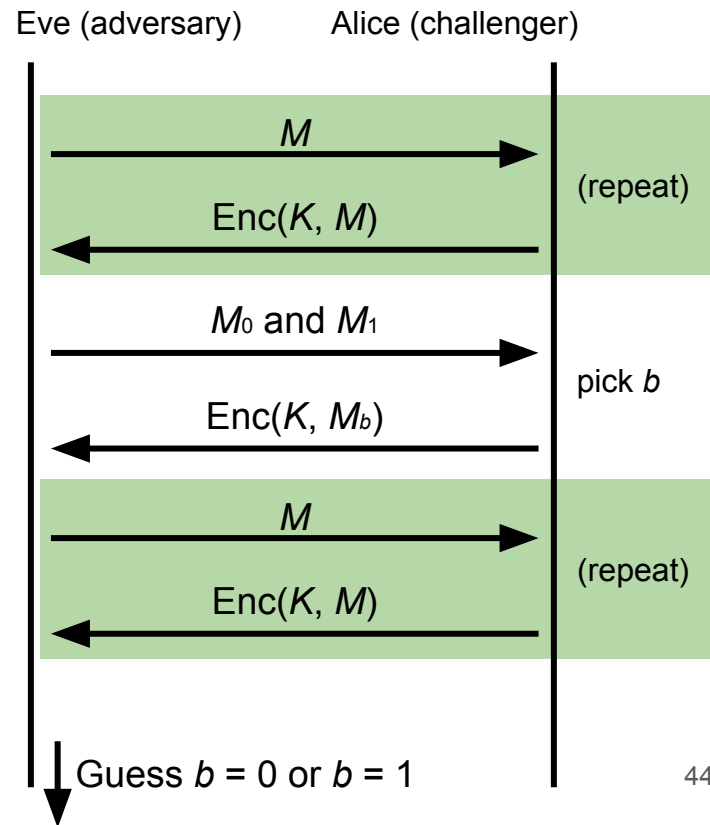
# One-Time Pads: Insecurity under IND-CPA

1. Eve can skip the first query phase...
2. What pair of messages should Eve send to Alice in the challenge phase?
  - Send  $M_0 = 0^k$ ,  $M_1 \neq M_0$
  - $M_0$  is a string of  $k$  0's,  $M_1$  is anything else
3. Alice chooses  $M_b$  to encrypt and sends the message back. Eve receives  $C_b = M_b \oplus K$
4. What messages should Eve for Alice to encrypt in the query phase?
  - Send  $M = M_0$  and receive  $C = M_0 \oplus K$
5. How can Eve know which message Alice encrypted?
  - If  $C_b = C$ , then Alice encrypted  $M_0$ . Otherwise, Alice encrypted  $M_1$ !
  - Eve wins the IND-CPA game with probability 1!




# One-Time Pads: Insecurity under IND-CPA

- Result: One-time pad is **not** IND-CPA secure
  - Yes, it is **perfectly secure** under some models and **very insecure** under other models!
  - This is the same reason as why two-time pad is insecure



# Traffic Analysis & Side Channels



# Traffic Analysis & Side Channels


- **Traffic analysis:** Analyzing who is talking to whom and when
  - The encryption schemes we'll be studying do not hide the identity of who you're talking to
  - The information used for this analysis is often referred to as **metadata**: Data *about* the message and its context
- **Side channels:** Information about the plaintext revealed as a result of the *implementation* of the scheme, not the scheme itself
  - Modern crypto systems are usually broken through side channels

# Traffic Analysis & Side Channels in Practice: Spies

Computer Science 161

Nicholas Weaver

- In the 1990s, there were some Russian spies in the US
  - The TV series “The Americans” was based on this incident
- A Cuban number station had a bug: some nights it never broadcasted “9”
  - Normally, 0–9 would be equally frequent
- It turns out this corresponded to when the Russian spies were on vacation
  - The way that random numbers were generated for cover traffic had a bug in it
  - The FBI used this as part of their investigation
- **Takeaway:** Secure algorithms can be broken in insecure implementations, leaking information



# Summary: IND-CPA and One-Time Pads

- IND-CPA security
  - Even if Eve can trick Alice into encrypting some messages of Eve's choosing, given the encryption of either  $M_0$  or  $M_1$ , Eve cannot distinguish which message was sent with probability greater than  $1/2$
  - We can use the IND-CPA game to test for IND-CPA security
  - Edge cases:
    - IND-CPA secure schemes can leak length
    - Eve is limited to polynomial-time algorithms, and must have a non-negligible advantage to win
- One-time pads
  - Symmetric encryption scheme: Alice and Bob share a secret key
  - Encryption and decryption: Bitwise XOR with the key
  - Can be perfectly secure, but also IND-CPA insecure
    - No information leakage if the key is never reused
    - Information leaks if the key is reused
  - Impractical for real-world usage, unless you're a spy
- Side channels: Information can be leaked due to *implementation* issues



# Block Ciphers

Textbook Chapter 6.4 & 6.5



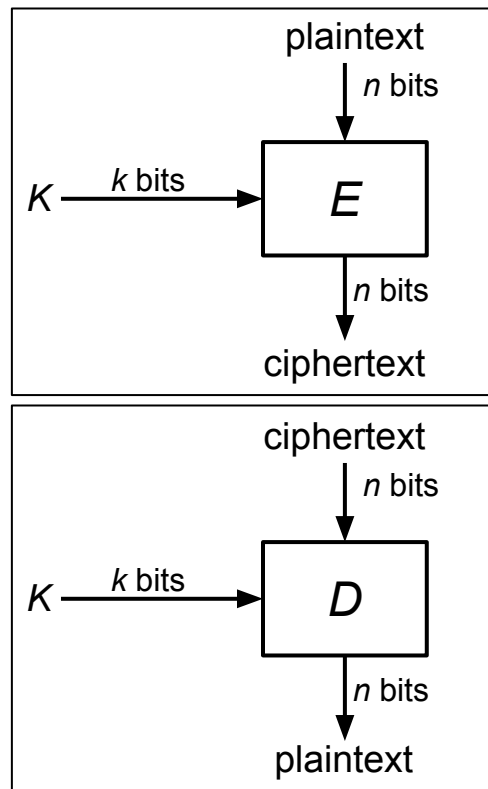
# Cryptography Roadmap

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none"><li>● One-time pads</li><li>● Block ciphers with chaining modes (e.g. AES-CBC)</li><li>● Stream ciphers</li></ul>	<ul style="list-style-type: none"><li>● RSA encryption</li><li>● ElGamal encryption</li></ul>
Integrity, Authentication	<ul style="list-style-type: none"><li>● MACs (e.g. HMAC)</li></ul>	<ul style="list-style-type: none"><li>● Digital signatures (e.g. RSA signatures)</li></ul>

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management


# Block Ciphers: Definition

- **Block cipher:** An encryption/decryption algorithm that encrypts a fixed-sized block of bits
- $E_K(M) \rightarrow C$ : Encryption
  - Inputs:  $k$ -bit key  $K$  and an  $n$ -bit plaintext  $M$
  - Output: An  $n$ -bit ciphertext  $C$
  - Sometimes written as:  $\{0, 1\}^k \times \{0, 1\}^n \rightarrow \{0, 1\}^n$
- $D_K(C) \rightarrow M$ : Decryption
  - Inputs: a  $k$ -bit key, and an  $n$ -bit ciphertext  $C$
  - Output: An  $n$ -bit plaintext
  - Sometimes written as:  $\{0, 1\}^k \times \{0, 1\}^n \rightarrow \{0, 1\}^n$
  - The inverse of the encryption function
- **Properties**
  - **Correctness:**  $E_K$  is a permutation,  $D_K$  is its inverse
  - **Efficiency:** Encryption/decryption should be fast
  - **Security:**  $E$  behaves like a random permutation




# Block Ciphers: Correctness

- $E_K(M)$  must be a **permutation (bijective function)** on  $n$ -bit strings
  - Each input must correspond to exactly one unique output
- Intuition
  - Suppose  $E_K(M)$  is not bijective
  - Then two inputs might correspond to the same output:  $E(K, x_1) = E(K, x_2) = y$
  - Given ciphertext  $y$ , you can't uniquely decrypt.  $D(K, y) = x_1$ ?  $D(K, y) = x_2$ ?





Not bijective: Two inputs encrypt to the same output



Bijective: Each input maps to exactly one unique output

# Block Ciphers: Security

- A secure block cipher behaves like a randomly chosen permutation from the set of all permutations on  $n$ -bit strings
  - A random permutation: Each  $n$ -bit input is mapped to one randomly-chosen  $n$ -bit output
- Defined by a distinguishing game
  - Eve gets two boxes: One is a randomly chosen permutation, and one is  $E_K$  with a randomly chosen key  $K$
  - Eve should not be able to tell which is which with probability  $> 1/2$



One of these is  $E_K$  with a randomly chosen  $K$ , and the other one is a randomly chosen permutation. Eve can't distinguish them.

# Block ciphers: Brute-force attacks?

- How hard is it to run a brute-force attack on a 128-bit key?
  - We have to try  $2^{128}$  possibilities. How big is  $2^{128}$ ?
- Handy approximation:  $2^{10} \approx 10^3$ 
  - $2^{128} = 2^{10 \cdot 12.8} \approx (10^3)^{12.8} \approx (10^3)^{13} = 10^{39}$
- Suppose we have massive hardware that can try  $10^9$  (1 billion) keys in 1 nanosecond (a billionth of a second). That's  $10^{18}$  keys per second
  - We'll need  $10^{39} / 10^{18} = 10^{21}$  seconds. How long is that?
  - One year  $\approx 3 \times 10^7$  seconds
  - $10^{21}$  seconds /  $3 \times 10^7 \approx 3 \times 10^{13}$  years  $\approx 30$  trillion years
- **Takeaway:** Brute-forcing a 128-bit key takes astronomically long. Don't even try.

# Block ciphers: Brute-force attacks?

- How hard is it to run a brute-force attack on a 256-bit key in the same time?
  - We need  $10^{52}$  of the brute-force devices from before
  - If each brute-force device from before is 1 cubic millimeter, this would take  $10^{43}$  cubic meters of space
  - That's the volume of  $7 \times 10^{15}$  suns!
  - For reference, the Milky Way galaxy has just  $10^{11}$  stars
- **Takeaway:** Brute-force attacks on modern block ciphers are not possible, assuming the key is random and secret
  - 128-bit key? Definitely not happening.
  - 256-bit key? Lol no.

# Block Ciphers: Efficiency

- Encryption and decryption should be computable in microseconds
  - Formally: `KeyGen()`, `Enc()`, and `Dec()`, should not take exponential time
- Block cipher algorithms typically use operations like XOR, bit-shifting, and small table lookups
  - Very fast on modern processors
- Modern CPUs provide dedicated hardware support for block ciphers



# DES (Data Encryption Standard)

- Designed in late 1970s
- Block size 64 bits ( $n = 64$ )
- Key size 56 bits ( $k = 56$ )
- NSA influenced two facets of its design
  - Altered some subtle internal workings in a mysterious way
  - Reduced key size from 64 bits to 56 bits
  - Made brute force attacks feasible for an attacker with massive computational resources (by 1970s standards)
- The algorithm remains essentially unbroken 40 years later
  - The NSA's tweaking hardened it against an attack publicly revealed a decade later
- However, modern computer speeds make it completely unsafe due to small key size
  - $\sim 6.4 \times 10^{16}$ , say  $10^{10}$  tries per second on my single desktop computer's Nvidia graphics card:  
Takes  $\sim 6.4 \times 10^6$  seconds or  $\sim 70$  days

# AES (Advanced Encryption Standard)


- 1997–2000: NIST (National Institute of Standards and Technology) in the US held a competition to pick a new block cipher standard
  - One of the finalists, Twofish, was designed by Berkeley professor and occasional CS 161 instructor David Wagner!
- Out of the 5 finalists:
  - Rijndael, Twofish, and Serpent had really good performance
  - RC6 had okay performance
  - Mars had ugly performance
- On any given computing platform, Rijndael was *never* the fastest
- But on every computing platform, Rijndael was *always* the second-fastest
  - Twofish and Serpent each had at least one compute platform they were bad at
- Rijndael was selected as the new block cipher standard

# AES (Advanced Encryption Standard)

- Key size 128, 192, or 256 bits ( $k = 128, 192, \text{ or } 256$ )
  - Actual cipher names are AES-128, AES-192, and AES-256
  - Paranoid people like the NSA use AES-256 keys, but AES-128 is just fine in practice
- Block size 128 bits ( $n = 128$ )
  - Note: The block size is still always 128 bits, regardless of key size
- You don't need to know how AES works, but you do need to know its parameters
  - [here's a comic](#)


# AES Algorithm

- Different key sizes use different numbers of rounds
  - 10 rounds for 128-bit keys
  - 12 rounds for 192-bit keys
  - 14 rounds for 256-bit keys
- Each round uses its own “round key” derived from the cipher key
- Each round:
  - SubBytes()
  - ShiftRows()
  - MixColumns() (if not last round)
  - AddRoundKey()




# AES Algorithm: SubBytes()

- Replace each byte in the block with another byte using an 8-bit substitution box




# AES Algorithm: ShiftRows()

- Cyclically shifts the bytes in each row by a certain offset
- The number of places each byte is shifted differs for each row




# AES Algorithm: MixColumns()

- Treats the 16-byte block as a  $4 \times 4$  matrix and multiply it by another matrix



# AES Algorithm: AddRoundKey()

- XOR the 16-byte block with the 16-byte round key



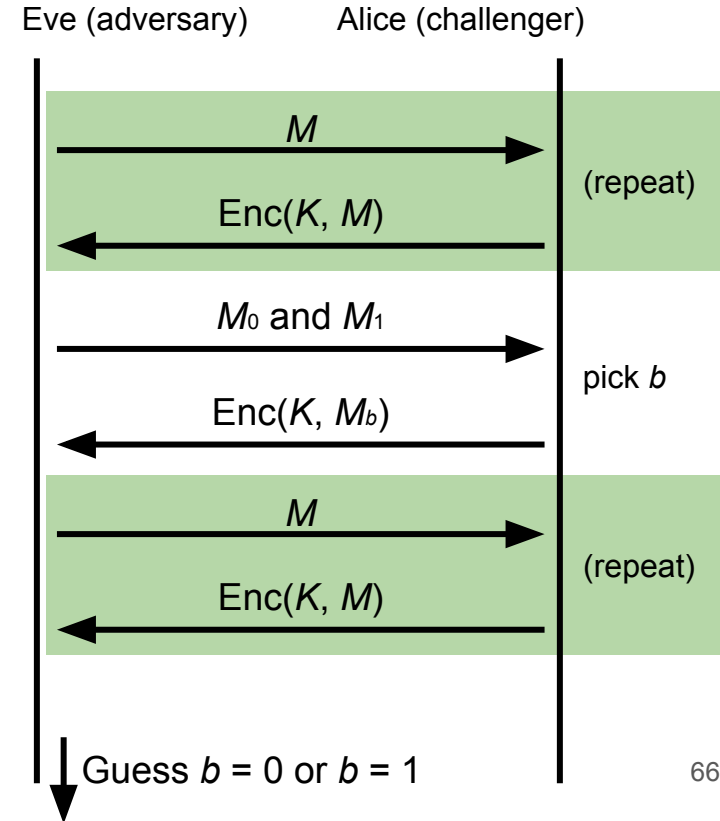


# AES (Advanced Encryption Standard)

- There is no formal proof that AES is secure (indistinguishable from a random permutation)
- However, in 20 years, nobody has been able to break it, so it is *assumed* to be secure
  - The NSA uses AES-256 for secrets they want to keep secure for the 40 years (even in the face of unknown breakthroughs in research)
- **Takeaway:** AES is the modern standard block cipher algorithm
  - The standard key size (128 bits) is large enough to prevent brute-force attacks

# Are Block Ciphers IND-CPA Secure?

- Consider the following adversary:
  - Eve sends two different messages  $M_0$  and  $M_1$
  - Eve receives either  $E_K(M_0)$  or  $E_K(M_1)$
  - Eve requests the encryption of  $M_0$  again
  - Strategy: If the encryption of  $M_0$  matches what she received, guess  $b = 0$ . Else, guess  $b = 1$ .
- Eve can win the IND-CPA game with probability 1!
  - Block ciphers are not IND-CPA secure



# Issues with Block Ciphers

- Block ciphers are not IND-CPA secure, because they're deterministic
  - A scheme is **deterministic** if the same input always produces the same output
  - No deterministic scheme can be IND-CPA secure because the adversary can always tell if the same message was encrypted twice
- Block ciphers can only encrypt messages of a fixed size
  - For example, AES can only encrypt-decrypt 128-bit messages
  - What if we want to encrypt something longer than 128 bits?
- To address these problems, we'll add **modes of operation** that use block ciphers as a building block!

# Summary: Block Ciphers

- Encryption: input a  $k$ -bit key and  $n$ -bit plaintext, receive  $n$ -bit ciphertext
- Decryption: input a  $k$ -bit key and  $n$ -bit ciphertext, receive  $n$ -bit plaintext
- Correctness: when the key is fixed,  $E_k(M)$  should be bijective
- Security
  - Without the key,  $E_k(m)$  is computationally indistinguishable from a random permutation
  - Brute-force attacks take astronomically long and are not possible
- Efficiency: algorithms use XORs and bit-shifting (very fast)
- Implementation: AES is the modern standard
- Issues
  - Not IND-CPA secure because they're deterministic
  - Can only encrypt  $n$ -bit messages